What Does a Streaming Client See When an LLM Server Crashes?

A Fault-Injection Study of Client-Visible Token Anomalies in vLLM and SGLang

Rajshekar Medipally
Independent researcher, Richmond, VA
https://github.com/rmedipallycic/llm-inference-fault-tolerance

October 2026

Abstract

LLM inference servers stream generated tokens to clients over requests that last seconds to minutes. When a serving process fails mid-generation, the client is left holding a partial response, and the recovery it performs determines what it ultimately shows a user or hands to a downstream agent. We built an open-source fault-injection harness that SIGKILLs an unmodified OpenAI-compatible server at a chosen token position, restarts it, applies a client recovery strategy, and records every streamed chunk, and applied it with identical code to vLLM and SGLang serving Llama 3.1 8B Instruct on A10G GPUs (86 runs, 430 trials). In both engines' default configurations the failure-free output for a fixed prompt was not unique at temperature 0, for different reasons: in vLLM it depended deterministically on what the server had served before, through automatic prefix caching; in SGLang the first request after each start varied from start to start. Because a crashed server restarts cold, naive retry contradicted text the client had already displayed in 25 of 30 vLLM and 24 of 30 SGLang trials on the main prompt, and on 3 of 5 (vLLM) and 4 of 5 (SGLang) further prompts. Recompute, which resubmits the delivered text, never contradicted it but usually changed the final answer. Re-tokenization was ruled out as the cause. With each engine's deterministic mode enabled (vLLM batch invariance, SGLang deterministic inference), recompute reproduced the failure-free output exactly in every temperature-0 trial (35 of 35) and naive retry never contradicted the client. Recovery latency was dominated by server restart (107–195 s). All raw logs and the table-generating script are public.

1 Introduction

LLM serving has become stateful streaming infrastructure. A generation request holds a growing key-value (KV) cache on an accelerator and emits tokens to the client incrementally. Failure handling, however, is still largely stateless: if the serving process dies, the client or a load balancer resubmits the request and generation starts again.

That recovery has a cost that server-side metrics miss. A client that has already displayed 128 tokens may receive those tokens a second time, may receive a continuation that disagrees with what it already showed, or may lose text at the boundary. For a chat interface this is a visible glitch. For an agent that has already acted on a partial tool call or a partial plan, it is a correctness problem. Stream processing systems name this property *delivery semantics* and engineer exactly-once output deliberately [5, 2, 1]. LLM serving has no equivalent definition, and to our knowledge no measurement of what clients actually receive.

A common intuition is that at temperature 0 a retry is harmless: greedy decoding is deterministic, so resubmitting the prompt reproduces the same text and the client only has to drop the duplicate. Our measurements on vLLM refute that intuition, and so do our measurements on SGLang, for a different reason. In vLLM the output for a fixed prompt depended on what the server had served earlier; in SGLang the first request after each server start varied from start to start. Either way, the restarted server that answers a retry does not reproduce the text the client was already shown.

This paper makes three contributions:

1. **Definitions.** Token-level delivery semantics for streamed generation, and client-visible anomalies defined against *two* failure-free references rather than one (Section 3).
2. **A fault-injection harness** for any OpenAI-compatible server, which kills the full server process tree at a chosen chunk, restarts it, applies a client recovery strategy, and records every chunk with timestamps, the code commit, package versions, and GPU inventory (Section 4).
3. **A measurement study** of naive retry and recompute on vLLM and SGLang at temperature 0 and 0.8 and three kill positions, with a cache ablation for each engine (Sections 5–6), which finds that in both engines' default configurations failure-free output depends on server state, identifies the

mechanism in vLLM, and quantifies what that does to each recovery strategy.

The study deliberately uses unmodified engines and client-side recovery only. It establishes the baseline that a server-side exactly-once protocol, the next stage of this project, would have to improve on.

2 Background

Autoregressive serving. A request runs a prefill phase, which processes the prompt in parallel and builds the KV cache, followed by a decode phase that produces one token per step and streams it to the client. Engines such as vLLM [7] and SGLang [14] batch requests at iteration granularity [13] and manage the KV cache in fixed-size blocks.

The state lost on failure. The KV cache is a deterministic function of the weights and the token sequence, so it can always be rebuilt by prefill. What cannot be rebuilt from the prompt alone is the *sampling path*: at temperature above zero, the continuation depends on random draws whose state lives in the dead process.

Rollback-recovery vocabulary. Classic rollback-recovery distinguishes checkpointing from logging and defines *output commit*: a message to the outside world must not be released until the state that produced it is recoverable [5]. Mapped onto LLM serving, the token sequence is a replayable input log, the KV cache is derived operator state, and the streaming client is the external sink where delivery semantics are decided.

3 Delivery model and anomalies

A failure occurs after the client has received delivered text D. After recovery the client receives recovery text R. What the client finally holds depends on the strategy: **naive retry** resubmits the original prompt with the original `max_tokens`, as many client retry wrappers do, and the retry is a complete answer, so the final text is $F = R$; **recompute** resubmits the prompt concatenated with D and requests the remaining token budget, so the server rebuilds the KV cache by prefill and continues, and $F = D \cdot R$. Neither requires engine changes.

Two references. Before the trials of each run, the harness issues the request twice without failures and records both outputs: G, the first request after the server starts, and G', the second. A reproducible server has $G = G'$. We do not assume this; we measure it, and compare every final text with both.

Anomalies.

- **Contradiction** (naive retry): R does not start with D, i.e. the replay disagrees with text the client has already displayed. A client that deduplicates by dropping the first $|D|$ characters of R would then splice two different texts. Recompute cannot contradict D by construction.
- **Duplication** (naive retry): the client receives D again. Without contradiction this is full duplication, which an idempotent client can remove.
- **Divergence**: F equals neither G nor G', i.e. the failure changed the final answer relative to any failure-free run we observed.
- **Recovery latency**: wall time from the kill to the first chunk delivered after recovery, which includes server restart.

Comparison is on text, not token ids, because the OpenAI-compatible streaming API does not guarantee one token per chunk.

4 Fault-injection harness

The harness (`harness/`, about 340 lines of Python) has three parts.

Server lifecycle. The server command is started in its own process group and declared ready when `GET /v1/models` returns 200. A fault is a `SIGKILL` of the whole process group, since vLLM and SGLang spawn worker processes that otherwise keep holding GPU memory and the port. On Windows, used only for laptop smoke tests, the equivalent is `taskkill /F /T`.

Streaming client. Requests go to `/v1/completions` with `stream=true`. Every server-sent event with non-empty text is recorded as a chunk with wall-clock and monotonic timestamps. The kill fires from the chunk callback when the configured chunk count is reached, and the client keeps reading, so chunks the server had already sent before dying are recorded as in-flight rather than lost.

Trial protocol. Each run (one engine, strategy, temperature, and kill position) starts the server, records two failure-free references G and G', then for each trial: streams the request, kills at the kill point, restarts the server and waits for readiness, issues the recovery request, and writes a JSON record containing every chunk, the kill time, the restart time, and the client view V.

Provenance. Each run records the git commit and whether the working tree was dirty, the full command line, installed versions of `vllm`, `sglang`, `torch`, and `transformers`, the GPU model, memory, and driver from `nvidia-smi`, and the full server log. Measured runs are written to `experiments/` and committed; harness smoke tests are written to a separate, git-ignored directory and are never reported. Every table in this paper is generated by one script, `analysis/make_tables.py`, from those committed files.

Validation. The harness was exercised end to end on a laptop CPU with llama.cpp and a 0.5B model before any GPU run. That check, documented in the repository as a mechanics check rather than a result, led to two changes: a duplication metric based on shared prefix fraction rather than any shared character, and the reference-reproducibility check, after two clean seeded runs on that server disagreed at temperature 0.8.

5 Methodology

Setup. Two identical AWS `g5.xlarge` instances (one NVIDIA A10G, 24 GB; 4 vCPUs each; same machine image and driver), one per engine, Llama 3.1 8B Instruct in bf16. The harness code was identical for both engines (commits `90da903` and `ddd8890` differ only outside `harness/`).

vLLM ran in its own Python 3.12 environment with default engine flags except a 4,096-token context cap (Llama 3.1's 128k default does not fit a 24 GB KV cache; every request here is far shorter). Defaults include automatic prefix caching, which the server logs confirm was on (`enable_prefix_caching=True`). The machine image ships no CUDA toolkit, so FlashInfer's sampler, which compiles at startup, was replaced by PyTorch sampling (`VLLM_USE_FLASHINFER_SAMPLER=0`); this is recorded in every run log.

SGLang could not be started directly on that image, because several of its components compile GPU code at startup (the attempts are documented in the repository). It ran instead from the official container image `lmsysorg/sglang:v0.5.21` (digest `sha256:b1259f3e...`), which includes a CUDA 13.0 toolkit, with the harness running inside the same container so that the process-group kill reaches the server. Flags were defaults except the same 4,096-token context cap and `--mem-fraction-static 0.85` (the default left no room for the KV cache on 24 GB). The server logs record what the defaults selected: FlashInfer attention and sampling, CUDA graphs, radix cache on, and deterministic inference off. The two engines therefore differ in kernels and sampler as well as in design; we compare each engine's default behaviour, not individual kernels. Versions are in Table 1.

86 measured runs, 430 trials, harness commit(s) `90da903, 9510da0, ddd8890`. In-flight chunks received after the kill: 0 at most, over all trials.

Grid. One fixed prompt (an explanation of write-ahead logging), `max_tokens`=256, per-request seed 1234. Strategy $\in$ {naive retry, recompute} $\times$ temperature $\in \{0, 0.8\}$ $\times$ kill point $\in$ {25, 128, 230} chunks (about 10%, 50%, and 90% of the budget), 5 trials each: 12 runs and 60 trials per engine, produced by `TRIALS=5 scripts/run_matrix.sh vllm` and by `scripts/sgl_grid.sh` (the same grid with the SGLang flags above). Requests are issued one at a time, so there is never more than one request in a batch.

Ablation. After the vLLM grid showed $G \neq G'$, we repeated the kill-at-128 configurations for both strategies and temperatures with automatic prefix caching disabled (`--no-enable-prefix-caching`), all else equal: 4 runs, 20 trials, labeled "no APC". For SGLang we ran the same four configurations with its radix cache, the corresponding feature, disabled (`--disable-radix-cache`), labeled "no radix".

Follow-up runs. After the first results, a second set of runs (harness commit `9510da0`, scripts `scripts/extras_*.sh`), all at the 128-chunk kill point with 5 trials each, tested the open questions: (i) each engine's deterministic mode, vLLM's batch invariance (`VLLM_BATCH_INVARIANT=1`) and SGLang's deterministic inference (`--enable-deterministic-inference`, with `sampling_seed` 1234 in each request), on the original prompt at both temperatures; (ii) SGLang with a fixed server seed (`--random-seed 1234`); (iii) a vLLM recompute variant, *recompute (ids)*, that resubmits token ids (the prompt's ids from `/tokenize` plus the delivered tokens' ids) instead of text, and records whether re-tokenizing the delivered text would have produced the same ids; and (iv) five further prompts (a TCP explanation, a short story, a coffee recipe, processes versus threads, and the 2008 financial crisis; `harness/prompts.json`) at temperature 0, for each engine in its default and its fixed configuration (vLLM without prefix caching, SGLang deterministic). That is 54 runs and 270 trials.

Hypotheses. Stated before the runs:

H1 At temperature 0, naive retry yields full duplication in every trial and recompute yields $F = G$ in every trial.

H2 At temperature 0.8 with a fixed seed, vLLM reproduces G across clean runs; if so, naive retry still yields full duplication and recompute diverges after the kill point.

Table 1: Software and hardware, from `run.json`.

Config	Engine	torch	GPU
vLLM	0.31.0	2.13.0	NVIDIA A10G (23028 MiB), driver 595.91.07
vLLM, no APC	0.31.0	2.13.0	NVIDIA A10G (23028 MiB), driver 595.91.07
vLLM, batch-inv.	0.31.0	2.13.0	NVIDIA A10G (23028 MiB), driver 595.91.07
SGLang	0.5.21	2.13.0+cu130	NVIDIA A10G (23028 MiB), driver 595.91.07
SGLang, no radix	0.5.21	2.13.0+cu130	NVIDIA A10G (23028 MiB), driver 595.91.07
SGLang, fixed seed	0.5.21	2.13.0+cu130	NVIDIA A10G (23028 MiB), driver 595.91.07
SGLang, determ.	0.5.21	2.13.0+cu130	NVIDIA A10G (23028 MiB), driver 595.91.07

H3 Recovery latency is dominated by server restart, not by prefill of the delivered text.

6 Results

vLLM: failure-free output was not unique. In all 12 runs of the default configuration, $G \neq G'$ (Table 2). At temperature 0 the two references agreed for the first 506 characters and then took different wording ("The log is scanned from the end…" versus "The log is scanned to determine…"). With a fixed seed at temperature 0.8 they diverged at character 9. Neither is noise: every run's G was identical to every other run's G at the same temperature, and likewise for G'. The output was a deterministic function of the server's history, not of the prompt alone. Within a run, all trials after the first produced identical text, so the 5 trials per configuration are repeated observations of at most two distinct server histories, not independent samples.

Table 2: Failure-free reproducibility: first request after server start (G) versus second (G'). LCP: their common prefix in characters (range over runs). Distinct G, G': how many different texts occurred across the runs of a row. "No APC", "no radix": the engine's prefix cache disabled.

Config	T	$G{=}G'$	LCP	$\#G$	$\#G'$
vLLM	0	0/6	506	1	1
vLLM	0.8	0/6	9	1	1
vLLM, no APC	0	2/2	1228	1	1
vLLM, no APC	0.8	2/2	1219	1	1
vLLM, batch-inv.	0	2/2	1228	1	1
vLLM, batch-inv.	0.8	2/2	1225	1	1
SGLang	0	0/6	267–676	4	1
SGLang	0.8	0/6	1–2	6	6
SGLang, no radix	0	0/2	267–412	2	1
SGLang, no radix	0.8	0/2	0–1	2	2
SGLang, fixed seed	0	0/2	267–554	2	2
SGLang, fixed seed	0.8	0/2	1	2	2
SGLang, determ.	0	2/2	1242	1	1
SGLang, determ.	0.8	0/2	3	1	1

vLLM: naive retry contradicted the client. The restarted server is always in the state that produced G, and the replay equaled G in all 30 naive-retry trials (Table 3). The text the client had already shown, D, came from a warm server and followed G'. Wherever the kill fell after the point at which G and G' part, the replay therefore contradicted D: in 25 of 30 trials, at character 506 at temperature 0 and at character 9 at temperature 0.8. The only clean case was temperature 0 with the kill at 25 chunks (113 characters, before character 506), where the replay began with D and the anomaly was pure, removable duplication. A client that drops the first $|D|$ characters of a replay, the usual deduplication, would in the other 25 trials have produced a spliced text that no failure-free run generated.

vLLM: recompute kept the shown text but changed the answer. Recompute never contradicted D, but its final text matched G in 4 of 30 trials and G' in 2 of 30. At temperature 0 with the kill at 25 chunks, all five final texts departed from both references at character 267, after the kill (113) and before G and G' part (506): a third greedy output. The same third point of departure appeared without any failure. In each recompute run, trials 1–4 ran on a server whose only earlier request was trial 0's recovery (prompt plus D), and their pre-kill streams, which no failure had yet touched, departed from both references in 16 of 24 trials. The two $F{=}G'$ cases are trial 0 at the 230-chunk kill; the four $F{=}G$ cases are trials 1–4 at temperature 0.8 with the kill at 25 chunks, where a 25-token recovery prefix left the continuation on the cold-server path.

vLLM hypotheses. **H1 is refuted**: at temperature 0, naive retry gave full duplication only at the earliest kill point (5 of 15 trials) and recompute never reproduced G (0 of 15). **H2's premise is refuted**: a fixed seed did not reproduce G across requests, so its predictions cannot be tested as stated. **H3 is supported** (Table 4): restart took 108–195 s and the first recovered token arrived 0.23–0.32 s after the server was ready. Recompute's extra prefill was visible but small: median time to first token after restart rose from 0.24 to 0.27 s at temperature 0 and from 0.28 to 0.31 s at 0.8 as the

kill moved from 25 to 230 chunks, against 0.23–0.28 s for naive retry. The higher median restart for recompute is not a strategy effect: restart happens before the recovery request is sent, and both initial startup and restart rose from about 110–122 s to 160–195 s from the third trial of the temperature-0, 230-chunk recompute run onward and stayed there for the rest of the grid; the cause was not identified. In the ablation, run later the same day, restarts were back to 107–120 s.

vLLM prefix-cache ablation. With automatic prefix caching disabled (rows "no APC"), $G = G'$ in all 4 runs at both temperatures (Table 2), and naive retry no longer contradicted the client: in 10 of 10 trials the replay equaled both references and began with D, so the anomaly was pure duplication that a client can remove (Table 3). The cache setting also changed the reference itself at temperature 0: the no-APC output departs from the default configuration's G at character 267, although G is the first request after start, when the cache is empty. At temperature 0.8 the no-APC reference equals the default G exactly.

Recompute behaved differently by temperature. At 0.8 its final text equaled G in 5 of 5 trials. At temperature 0 it equaled neither reference in 5 of 5: all five departed at character 655, 43 characters after the kill (612), so recompute changed the greedy answer even with no prefix cache. That final text is character-for-character identical to the default configuration's recompute text at the 25-chunk kill. Recompute's answer at temperature 0 is therefore not explained by the cache. It is consistent with the delivered text being processed as prompt (prefill) rather than generated (decode), or with re-tokenization of D; these data do not separate the two.

SGLang: the first request after a start was not reproducible. SGLang showed a different pattern (Table 2). At temperature 0 the second request on a running server, G', was the same text in all 6 runs, but the first request after a start was not: the 21 first-after-start outputs we recorded (6 references G and 15 naive-retry replays) contained 10 distinct texts, none equal to G', departing from it at characters 267 to 676. Disabling the radix cache did not change this: G' was again the same in both runs, and 7 first-after-start outputs gave 4 distinct texts. In SGLang the output therefore depended on whether the server had just started, not on cache contents, and the cold output was not repeatable from one start to the next. We did not identify the source. One untested candidate is that, with no seed given, SGLang chooses a new server random seed at every start; the logs record a different value each time.

At temperature 0.8 the request seed had no observable effect: G and G' diverged within the first two characters in all 8 runs, and the 20 naive-retry pre-kill streams were all different. SGLang provides seeded, reproducible sampling through an opt-in deterministic mode with a seed in its own sampling parameters [9]; our harness sends the OpenAI-style `seed` and did not enable that mode, so SGLang's 0.8 results describe unseeded sampling.

SGLang: the same consequences. Naive retry contradicted text already shown in 24 of 30 default-configuration trials: at temperature 0 in 0, 4 and 5 of 5 trials as the kill moved from 25 to 128 to 230 chunks (the warm stream D followed G', the cold replay did not), and at 0.8 in 15 of 15. Without the radix cache it did so in 10 of 10. Recompute never contradicted D, and its final text matched neither failure-free output of its run in 26 of 30 default-configuration trials. Recovery was again dominated by restart (Table 4): 125–145 s with the radix cache on, and the first recovered token 0.33–0.51 s after the server was ready. Restarts in the no-radix runs, which ran last, took 172–195 s; as with vLLM's drift, we did not identify the cause and do not attribute it to the cache setting. For SGLang, as for vLLM, **H1 is refuted** (at temperature 0 naive retry gave full duplication only at the earliest kill point, and recompute reproduced G in 3 of 15 trials), **H2's premise is refuted** (the seed had no observable effect), and **H3 is supported**.

7 Follow-up experiments

Re-tokenization is not the cause. In all 10 recompute (ids) trials, re-tokenizing the delivered text gave exactly the token ids the engine had generated (Table 5). Resubmitting the ids instead of the text changed nothing: without prefix caching, the final text again departed from the reference at character 655, the same place as text-based recompute. The remaining explanation for recompute's change of answer is that the delivered tokens' KV cache, computed in one prefill on recovery, differs numerically from the one built token by token during decoding. Requesting token ids on every call also changed vLLM's default-configuration reference to the no-APC text, so the default-configuration ids run serves only as a re-tokenization check.

SGLang's server seed is not the cause. With `--random-seed 1234`, SGLang's first request after a start still varied from start to start (Table 2) and naive retry still contradicted the client in 5 of 5 trials at temperature 0.

Deterministic modes made recompute exact. With vLLM's batch invariance, $G = G'$ in all 4 runs, naive retry never contradicted the client (0 of 10), and

recompute reproduced the failure-free output exactly in 10 of 10 trials, at temperature 0 and at 0.8 with a fixed seed (Table 3). Its temperature-0 output was identical to the no-APC output. With SGLang's deterministic inference at temperature 0, $G = G'$, naive retry never contradicted the client, and recompute was exact in 5 of 5. That is consistent with the explanation above: kernels that give the same result regardless of batch and prefill shape remove the difference between the recovered and the original computation. At temperature 0.8, SGLang's deterministic mode made each reference repeatable across runs, but G and G' still diverged at character 3; naive retry contradicted the client in 5 of 5 trials and recompute was not exact (0 of 5), so its sampling was not reproducible across requests in this configuration.

The pattern holds across prompts. Table 6 repeats the 128-chunk, temperature-0 configuration on five more prompts. In vLLM's default configuration, $G \neq G'$ for three of the five prompts, and for exactly those three naive retry contradicted the client in 5 of 5 trials (15 of 25 overall); for the other two, the first and second requests happened to agree. Without prefix caching, $G = G'$ for all five and naive retry never contradicted (0 of 25), but recompute was exact for only one prompt (5 of 25). In SGLang's default configuration $G = G'$ in only 1 of 10 runs, naive retry contradicted the client in 14 of 25 trials, and recompute was exact in 3 of 25. In SGLang's deterministic mode, all 10 runs had $G = G'$, naive retry never contradicted (0 of 25), and recompute was exact in 25 of 25.

Cost. The deterministic modes did not slow recovery much in these runs: restart times were in the same range, and the first token after restart arrived in 0.26–0.30 s with vLLM's batch invariance. With SGLang's deterministic mode it arrived in 0.34–0.46 s in some runs and 3–4 s in others, which we did not investigate. Throughput under load, which is what batch-invariant kernels are reported to cost [6], was not measured.

8 Discussion

Two engines, two mechanisms, one consequence. vLLM's variation was deterministic and tied to its prefix cache; SGLang's was tied to server start and was not repeatable across starts. A client cannot tell these apart, and for a client they have the same effect: the server that answers a retry is freshly started, and a freshly started server did not reproduce the text the client had been shown, in 25 of 30 vLLM and 24 of 30 SGLang trials.

Retry is not replay. The usual argument for retrying an LLM request is that at temperature 0 it is idempotent up to duplication. On the default vLLM configuration measured here, it was not: the server that answers a retry is a fresh process, and a fresh process gave a different greedy answer from the warm one that had been streaming. For a chat client this means a visible rewrite of text already shown in most trials on both engines; for an agent that acted on a partial tool call or plan, it means the action was based on text that the final answer contradicts.

What a client and an operator should do. With default engine settings, recompute is the only client-side strategy that never contradicted what was shown, at the price of a final answer that no failure-free run produced, and a few tens of milliseconds of extra prefill. Naive retry with deduplication is unsafe unless the client verifies that the replay begins with D, and should replace rather than splice when it does not. On vLLM, an operator can remove the contradictions by disabling automatic prefix caching, as the ablation did; we did not measure the throughput cost, which on workloads with shared prefixes is the reason the cache exists. On SGLang, disabling the radix cache did not help. The combination that worked on both engines was an engine deterministic mode on the server and recompute on the client: at temperature 0 it reproduced the failure-free output exactly in all 35 trials (Section 7), which is what a client sees under exactly-once delivery, without any change to the engines' code.

Why client-side recovery is not enough. Even in the best case for recompute, the client cannot restore the server's state: neither the sampler's random state above temperature 0 nor, as measured here, whatever server state made G differ from G'; and even with that state removed by disabling the prefix cache, recompute at temperature 0 did not reproduce the original continuation. Deterministic modes close this gap at temperature 0, but not, in SGLang, for sampled output, and they depend on the operator's configuration. Exactly-once delivery in general therefore requires state the client does not hold: a sequence-numbered token log, and the sampler state, committed before tokens are released, with recovery continuing from the log rather than regenerating. That protocol is the next stage of this project, and these measurements are its baseline.

9 Related work

Fault-tolerant LLM serving. DéjàVu [10] streams and replicates KV cache to recover from failures and reports server-side recovery time; it does not characterize what streaming clients receive. SpotServe [8] mi-

grates in-flight context when preemptible instances are reclaimed with warning, and Llumnix [11] live-migrates requests for load balancing. Both handle planned moves; this study targets unplanned fail-stop crashes. **Serving engines.** Orca [13] introduced iteration-level scheduling, vLLM [7] block-managed KV memory, and SGLang [14] prefix sharing via RadixAttention. **Delivery semantics.** Chandy and Lamport [3] define consistent snapshots, Elnozahy et al. [5] output commit, and Flink [2] and Spark Structured Streaming [1] build exactly-once sinks on them. **Inference nondeterminism.** He et al. [6] show that LLM inference at temperature 0 is not reproducible across server load because kernels are not batch-invariant, and build batch-invariant kernels. Our observation is the fault-tolerance consequence: requests here were issued one at a time, yet output still depended on what the server had served before, so a retry after a crash is not a replay. Optimal checkpoint intervals [12, 4] motivate later stages of this project.

10 Threats to validity

- **Trials are not independent samples.** In vLLM, output was deterministic given server history, so 5 trials per configuration are repeated observations of one or two histories, and counts such as "25 of 30" describe how many configurations and histories were affected, not a sampled rate. In SGLang the cold output varied from start to start, so its trials do sample that variation, but 5 per configuration is too few to estimate rates.

- **Six prompts, one model, one GPU type.** The main grid used one prompt; five more were tested only at temperature 0 and the 128-chunk kill point, and vLLM's batch invariance only on the main prompt. Divergence positions are prompt-specific (two of five extra prompts showed no vLLM divergence within 256 tokens). Other models and tensor-parallel serving may differ.

- **Default configurations, different kernels.** Each engine ran with its own defaults on this GPU (vLLM with PyTorch sampling because FlashInfer's sampler could not compile on the host image; SGLang with FlashInfer attention and sampling inside its container). The comparison is between default behaviours, not kernels.

- **SGLang seed.** The harness's OpenAI-style seed had no observable effect on SGLang, so its temperature-0.8 results are for unseeded sampling and are not comparable to vLLM's seeded runs. Its temperature-0 results do not involve sampling.

- **vLLM sampler substitution.** vLLM ran with PyTorch sampling instead of FlashInfer's sampler, because of the missing CUDA toolkit. The temperature-0 results do not involve random sampling, but the 0.8 results may differ under FlashInfer.

- **Mechanism.** For vLLM, the ablation shows that the dependence on server history requires automatic prefix caching, and re-tokenization was ruled out for recompute. For SGLang, neither the radix cache nor the server seed was required, and the source of the start-to-start variation was not identified. That both engines' deterministic modes removed every contradiction and made recompute exact at temperature 0 supports, but does not prove, a numerical explanation [6]; we did not identify which kernel differs. The ablations and follow-up runs covered only the 128-chunk kill point.

- **Re-tokenization in recompute.** Recompute sends D as text, which the engine re-tokenizes. We checked this directly for vLLM (Section 7): re-tokenization matched in all 10 trials checked, at the 128-chunk kill on the main prompt. Other prompts or kill points could tokenize differently at the boundary.

- **Fail-stop only; latency includes restart.** `SIGKILL` is a clean crash; GPU hangs and silent corruption are not modeled. With a single instance, recovery includes model reload; a warm standby would be faster. Restart time also drifted during the grid (see H3).

11 Conclusion and future work

In the default configurations of both vLLM and SGLang, the failure-free output for a fixed prompt was not unique at temperature 0, so the restarted server that answers a retry did not reproduce what the client had been shown: naive retry contradicted displayed text in most trials on both engines, and on three (vLLM) and four (SGLang) of five further prompts. The cause was not re-tokenization; in vLLM it required prefix caching, and in SGLang it was tied to server start. Enabling each engine's deterministic mode and recovering by recompute reproduced the failure-free output exactly in every temperature-0 trial we ran. A retry is not a replay by default, but with the right server configuration and client strategy it can be, for greedy decoding.

Next stages: (0) the throughput cost of the deterministic modes, other models, and sampled output under SGLang's deterministic mode; (1) a gateway that assigns sequence numbers, logs tokens and sampler state

before release, and continues from the log on recovery, evaluated on the same harness for zero anomalies and per-token overhead; (2) measured recompute-versus-checkpoint cost curves for the KV cache; (3) an adaptive per-request checkpoint policy.

Artifact. Code, raw logs, and the table-generating script are at `https://github.com/rmedipallycic/llm-inference-fault-tolerance` (Apache-2.0). Every number above traces to a committed log file.

Acknowledgments

The harness code, analysis scripts, and much of the text of this paper were developed with substantial help from an AI assistant (Claude, Anthropic). The author directed the work, ran all experiments on AWS, and is responsible for every reported result; each number is generated by `analysis/make_tables.py` from the raw logs committed to the repository.

Table 3: Client-visible outcome per configuration (5 trials each), generated from the raw logs by `analysis/make_tables.py`. $|D|$: characters delivered before the kill (two values, or a range, when trials differed). Contradicts D: the naive replay disagreed with text already shown. $F{=}G$, $F{=}G'$: the final text equals the first or second failure-free reference.

Config	Strategy	T	Kill (chunks)	$\lvert D\rvert$ (chars)	Contradicts D	$F{=}G$	$F{=}G'$
vLLM	naive	0	25	113	0/5	5/5	0/5
vLLM	naive	0	128	643	5/5	5/5	0/5
vLLM	naive	0	230	1145	5/5	5/5	0/5
vLLM	naive	0.8	25	105	5/5	5/5	0/5
vLLM	naive	0.8	128	590	5/5	5/5	0/5
vLLM	naive	0.8	230	1079	5/5	5/5	0/5
vLLM	recomp.	0	25	113	n/a	0/5	0/5
vLLM	recomp.	0	128	625/643	n/a	0/5	0/5
vLLM	recomp.	0	230	1143/1145	n/a	0/5	1/5
vLLM	recomp.	0.8	25	105/114	n/a	4/5	0/5
vLLM	recomp.	0.8	128	556/590	n/a	0/5	0/5
vLLM	recomp.	0.8	230	1079/1085	n/a	0/5	1/5
vLLM	recomp. (ids)	0	128	625/626	n/a	0/5	1/5
vLLM, no APC	naive	0	128	612	0/5	5/5	5/5
vLLM, no APC	naive	0.8	128	601	0/5	5/5	5/5
vLLM, no APC	recomp.	0	128	612	n/a	0/5	0/5
vLLM, no APC	recomp.	0.8	128	601	n/a	5/5	5/5
vLLM, no APC	recomp. (ids)	0	128	612	n/a	0/5	0/5
vLLM, batch-inv.	naive	0	128	612	0/5	5/5	5/5
vLLM, batch-inv.	naive	0.8	128	556	0/5	5/5	5/5
vLLM, batch-inv.	recomp.	0	128	612	n/a	5/5	5/5
vLLM, batch-inv.	recomp.	0.8	128	556	n/a	5/5	5/5
SGLang	naive	0	25	113	0/5	1/5	0/5
SGLang	naive	0	128	643	4/5	1/5	0/5
SGLang	naive	0	230	1144	5/5	0/5	0/5
SGLang	naive	0.8	25	76–121	5/5	0/5	0/5
SGLang	naive	0.8	128	282–662	5/5	0/5	0/5
SGLang	naive	0.8	230	1120–1173	5/5	0/5	0/5
SGLang	recomp.	0	25	113	n/a	0/5	0/5
SGLang	recomp.	0	128	623/643	n/a	0/5	0/5
SGLang	recomp.	0	230	1113–1163	n/a	3/5	1/5
SGLang	recomp.	0.8	25	100–136	n/a	0/5	0/5
SGLang	recomp.	0.8	128	608–686	n/a	0/5	0/5
SGLang	recomp.	0.8	230	1115–1206	n/a	0/5	0/5
SGLang, no radix	naive	0	128	643	5/5	1/5	0/5
SGLang, no radix	naive	0.8	128	599–682	5/5	0/5	0/5
SGLang, no radix	recomp.	0	128	643	n/a	0/5	2/5
SGLang, no radix	recomp.	0.8	128	430–683	n/a	0/5	0/5
SGLang, fixed seed	naive	0	128	643	5/5	1/5	0/5
SGLang, fixed seed	naive	0.8	128	456–691	5/5	0/5	0/5
SGLang, fixed seed	recomp.	0	128	623/625	n/a	0/5	0/5
SGLang, fixed seed	recomp.	0.8	128	588–670	n/a	0/5	0/5
SGLang, determ.	naive	0	128	625	0/5	5/5	5/5
SGLang, determ.	naive	0.8	128	666	5/5	5/5	0/5
SGLang, determ.	recomp.	0	128	625	n/a	5/5	5/5
SGLang, determ.	recomp.	0.8	128	666	n/a	0/5	0/5

Table 4: Latency in seconds, median (min–max). Restart: kill to server ready. TTFT: server ready to first recovered chunk. Recovery: kill to first recovered chunk.

Config	Strategy	Trials	Restart	TTFT after restart	Recovery
vLLM	naive	30	116.6 (108.1–122.1)	0.26 (0.23–0.28)	116.9 (108.3–122.4)
vLLM	recomp.	30	177.7 (110.1–195.2)	0.27 (0.23–0.32)	178.0 (110.4–195.5)
vLLM	recomp. (ids)	5	116.1 (107.1–119.1)	0.59 (0.58–0.60)	116.7 (107.7–119.7)
vLLM, no APC	naive	10	109.6 (107.1–120.1)	0.25 (0.23–0.28)	109.9 (107.3–120.4)
vLLM, no APC	recomp.	10	111.1 (108.1–120.1)	0.28 (0.25–0.30)	111.4 (108.4–120.4)
vLLM, no APC	recomp. (ids)	5	116.1 (108.1–119.1)	0.61 (0.60–0.62)	116.7 (108.7–119.8)
vLLM, batch-inv.	naive	10	122.6 (119.1–131.1)	0.27 (0.26–0.28)	122.9 (119.4–131.4)
vLLM, batch-inv.	recomp.	10	128.1 (118.1–132.1)	0.29 (0.28–0.30)	128.4 (118.4–132.4)
SGLang	naive	30	133.2 (125.1–145.2)	0.40 (0.33–0.49)	133.6 (125.6–145.5)
SGLang	recomp.	30	129.2 (123.1–193.2)	0.42 (0.35–0.51)	129.6 (123.5–193.7)
SGLang, no radix	naive	10	184.7 (178.2–188.2)	0.40 (0.35–0.42)	185.1 (178.6–188.6)
SGLang, no radix	recomp.	10	185.7 (172.2–195.2)	0.43 (0.40–0.48)	186.1 (172.6–195.7)
SGLang, fixed seed	naive	10	132.7 (129.1–142.2)	0.41 (0.35–0.47)	133.1 (129.6–142.6)
SGLang, fixed seed	recomp.	10	131.6 (127.1–135.1)	0.44 (0.37–0.46)	132.1 (127.5–135.6)
SGLang, determ.	naive	10	133.6 (127.2–146.2)	1.72 (0.34–4.02)	135.4 (130.5–146.5)
SGLang, determ.	recomp.	10	134.6 (127.2–142.2)	1.70 (0.38–4.07)	135.1 (130.6–142.6)

Table 5: Recompute (ids), vLLM, temperature 0, kill at 128 chunks. Re-tok.: re-tokenizing D gave the generated token ids.

Config	T	n	Re-tok.	$F{=}G$	$F{=}G'$
vLLM	0	5	5/5	0/5	1/5
vLLM, no APC	0	5	5/5	0/5	0/5

Table 6: Five further prompts, temperature 0, kill at 128 chunks, 5 trials per strategy. Contra.: naive replay contradicted D. Exact: recompute's final text equals G or G'.

Config	Prompt	$G{=}G'$	Contra.	Exact
vLLM	coffee	0/2	5/5	0/5
vLLM	crisis	2/2	0/5	0/5
vLLM	story	0/2	5/5	0/5
vLLM	tcp	2/2	0/5	0/5
vLLM	threads	0/2	5/5	0/5
vLLM, no APC	coffee	2/2	0/5	5/5
vLLM, no APC	crisis	2/2	0/5	0/5
vLLM, no APC	story	2/2	0/5	0/5
vLLM, no APC	tcp	2/2	0/5	0/5
vLLM, no APC	threads	2/2	0/5	0/5
SGLang	coffee	0/2	3/5	0/5
SGLang	crisis	0/2	4/5	0/5
SGLang	story	0/2	5/5	0/5
SGLang	tcp	1/2	0/5	2/5
SGLang	threads	0/2	2/5	1/5
SGLang, determ.	coffee	2/2	0/5	5/5
SGLang, determ.	crisis	2/2	0/5	5/5
SGLang, determ.	story	2/2	0/5	5/5
SGLang, determ.	tcp	2/2	0/5	5/5
SGLang, determ.	threads	2/2	0/5	5/5

References

[1] Michael Armbrust, Tathagata Das, Joseph Torres, Burak Yavuz, Shixiong Zhu, Reynold Xin, Ali Ghodsi, Ion Stoica, and Matei Zaharia. Structured streaming: A declarative API for real-time applications in Apache Spark. In *Proceedings of the 2018 International Conference on Management of Data (SIGMOD)*, pages 601–613, 2018.

[2] Paris Carbone, Gyula Fóra, Stephan Ewen, Seif Haridi, and Kostas Tzoumas. Lightweight asynchronous snapshots for distributed dataflows, 2015. arXiv:1506.08603.

[3] K. Mani Chandy and Leslie Lamport. Distributed snapshots: Determining global states of distributed systems. *ACM Transactions on Computer Systems*, 3(1):63–75, 1985.

[4] J. T. Daly. A higher order estimate of the optimum checkpoint interval for restart dumps. *Future Generation Computer Systems*, 22(3):303–312, 2006.

[5] E. N. Elnozahy, Lorenzo Alvisi, Yi-Min Wang, and David B. Johnson. A survey of rollback-recovery protocols in message-passing systems. *ACM Computing Surveys*, 34(3):375–408, 2002.

[6] Horace He and Thinking Machines Lab. Defeating nondeterminism in LLM inference. Thinking Machines Lab: Connectionism, 2025. `https://thinkingmachines.ai/blog/defeating-nondeterminism-in-llm-inference/`.

[7] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with PagedAttention. In *Proceedings of the 29th Symposium on Operating Systems Principles (SOSP)*, 2023.

[8] Xupeng Miao, Chunan Shi, Jiangfei Duan, Xiaoli Xi, Dahua Lin, Bin Cui, and Zhihao Jia. SpotServe: Serving generative large language models on preemptible instances. In *Proceedings of ASPLOS*, pages 1112–1127, 2024.

[9] SGLang Team. Towards deterministic inference in SGLang and reproducible RL training. LMSYS Org blog, 2025. `https://lmsys.org/blog/2025-09-22-sglang-deterministic/`.

[10] Foteini Strati, Sara McAllister, Amar Phanishayee, Jakub Tarnawski, and Ana Klimovic. DéjàVu: KV-cache streaming for fast, fault-tolerant generative LLM serving. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, volume 235 of *PMLR*, pages 46745–46771, 2024.

[11] Biao Sun, Ziming Huang, Hanyu Zhao, Wencong Xiao, Xinyi Zhang, Yong Li, and Wei Lin. Llumnix: Dynamic scheduling for large language model serving. In *18th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*, 2024.

[12] John W. Young. A first order approximation to the optimum checkpoint interval. *Communications of the ACM*, 17(9):530–531, 1974.

[13] Gyeong-In Yu, Joo Seong Jeong, Geon-Woo Kim, Soojeong Kim, and Byung-Gon Chun. Orca: A distributed serving system for transformer-based generative models. In *16th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*, pages 521–538, 2022.

[14] Lianmin Zheng, Liangsheng Yin, Zhiqiang Xie, Chuyue Sun, Jeff Huang, Cody Hao Yu, Shiyi Cao, Christos Kozyrakis, Ion Stoica, Joseph E. Gonzalez, Clark Barrett, and Ying Sheng. SGLang: Efficient execution of structured language model programs. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. arXiv:2312.07104.